

Scheduling

Multi-Level Feedback Queue

- optimize turn around
- Responsive

has multiple queues that represent different priorities then works off Round Robin

Rule 1: $\text{Priority}(A) > \text{Priority}(B)$? A Runs

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, Round Robin

MLFQ tries to learn about processes as they run, and bases priority off that



Changing Priority

Rule 3: when a Job enters the system it is placed at the highest priority

Rule 4a: If Job uses entire time slice while running its priority is reduced

Rule 4b: If a Job gives up CPU before time runs out it keeps priority

Problems with MLFQ

Starvation: If there are too many high priority tasks, low priority jobs get starved

Gaming the Scheduler: A smart programmer can make their program take 99% and monopolize the system

Programs change Behavior: with current system it couldn't move back up the queue

* Priority Boost *

Rule 5: After some time "S" all jobs are moved to the topmost queue

* Better Accounting *

Revised Rule 4: Once a job uses up its time allotment at a given level (Regardless of how many times it has given up the CPU) its priority is reduced

- Defaults *

≈ 6 queues with slowly increasing time slices.

Priority boost every ≈ 1 second

Proportional - Share Scheduling

- Lottery: uses Randomness
- Stride: Deterministic
- CFS (Completely Fair Scheduler) = weighted Round Robin with dynamic time slices
 - uses Red/Black Trees
 - Scales well